

CA3

March 19, 2024

0.0.1 Imports

```
[1022]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

0.0.2 Reading data

```
[1023]: df= pd.read_csv("train.csv",index_col=0 )
```

0.0.3 Data exploration and visualisation

```
[1024]: # # for comparing different features to each other from the dataset, with hue
↳as Edible
# #I just made comments out of this so you dont have to read all the output

# sns.pairplot(df, hue='Edible')
# plt.show()
```

```
[1025]: #For looking at the distribution of the data of each feature in the dataset
#I just made comments out of this so you dont have to read all the output

# for column in df.columns:
#     plt.figure(figsize=(10, 4))
#     sns.histplot(df[column], kde=True) # kde (Kernel Density Estimate) adds
↳a density curve
#     plt.title(f'Distribution of {column}')
#     plt.xlabel(column)
#     plt.ylabel('Frequency')
#     plt.show()
```

```
[1026]: #Looking for correlation between features in the dataset
#I just made comments out of this so you dont have to read all the output
```

```
# corr_matrix = df.corr()
# # Apply formatting to round the values to two decimal places.
# # The .applymap() method applies a lambda function to each element of the
# ↪ correlation matrix.
# # The lambda function converts each value 'x' in the matrix to a formatted
# ↪ string with two decimal places.
# corr_matrix_rounded = corr_matrix.applymap(lambda x: f'{x:.2f}')

# # Plotting the heatmap, first setting the size of the heatmap
# plt.figure(figsize=(10, 8))
# # 'annot=corr_matrix_rounded' is used to display the rounded correlation
# ↪ values as annotations on the heatmap.
# # 'fmt=' specifies that the annotation texts (rounded values) are strings,
# ↪ not floats or integers.
# # 'cmap='coolwarm' sets the colormap to 'coolwarm'
# sns.heatmap(corr_matrix, annot=corr_matrix_rounded, fmt='', cmap='coolwarm')
# plt.title('Correlation Heatmap')
# plt.show()
```

```
[1027]: # Looking at df to see how many rows and columns there is
#I just made comments out of this so you dont have to read all the output
#df
```

```
[1028]: df.describe()
```

```
[1028]:
```

	Acoustic Firmness Index	Atmospheric Pressure at Harvest (Pa)	\
count	1245.000000	1248.000000	
mean	21.570077	101327.543269	
std	8.131888	4772.582203	
min	0.600000	83825.000000	
25%	17.300000	98095.750000	
50%	21.300000	101357.000000	
75%	25.300000	104470.750000	
max	156.520701	115636.000000	

	Bitterness Scale	Circumference (mm)	Color Intensity (a.u.)	\
count	1244.000000	1246.000000	1248.000000	
mean	1.808682	211.046062	63.300962	
std	1.960279	24.652278	7.118135	
min	0.000000	145.867667	46.060000	
25%	0.000000	192.982222	58.577500	
50%	1.000000	212.429730	62.390000	
75%	3.000000	226.313333	67.202500	
max	7.000000	274.719407	97.810000	

```
Find Distance from Main Vulcano (km) Length (mm) \
```

count	1246.000000	1244.000000
mean	44.188204	33.589646
std	28.331532	3.874355
min	2.550790	23.451799
25%	21.203453	30.777936
50%	38.880858	33.784401
75%	60.680816	36.049083
max	109.952511	43.691515

	Luminescence Intensity (a.u.)	Magnetic orientation (degree) \
count	1247.000000	1245.000000
mean	0.002850	178.933737
std	0.003607	102.113321
min	0.000004	0.085357
25%	0.000627	90.582930
50%	0.001757	180.043851
75%	0.003767	266.481765
max	0.044624	359.443812

	Odor index (a.u.)	Seed Count	Skin Thickness (mm) \
count	1246.000000	1247.000000	1247.000000
mean	16.602442	48.926812	3.839615
std	14.882420	35.871733	1.483029
min	1.722825	0.000000	1.000000
25%	7.615908	15.852008	3.000000
50%	12.258785	47.932723	4.000000
75%	20.302048	80.413165	5.000000
max	152.041780	112.968004	8.500000

	Soil pH where Grown	Sugar Content (mg)	Weight (mg)	pH \
count	1245.000000	1245.000000	1244.000000	1244.000000
mean	7.278739	11835.978313	78446.779224	4.614780
std	0.529815	2417.733775	26637.766132	0.993407
min	5.710000	4566.000000	32352.182600	-5.120000
25%	6.910000	10353.000000	57087.682600	4.174979
50%	7.280000	11807.000000	77440.682600	4.910000
75%	7.660000	13534.000000	92664.102600	5.080000
max	8.980000	18246.000000	164679.342600	6.081918

	Edible
count	1248.000000
mean	0.494391
std	0.500169
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000

max 1.000000

Comments on the data after visualization Some of the pH values didn't look right. There was some negative pH values, which doesn't make sense. It looked like there was some outliers on some of the features, which is also getting removed before training.

0.0.4 Data cleaning

```
[1029]: #replacing nan-values with median
# I do this instead of removing the rows, because it gave me better accuracy
↳when training models
df_cleaned= df.fillna(df.median())

# Dropping rows with negative pH values
df_cleaned = df_cleaned.drop(df_cleaned[df_cleaned['pH'] < 0].index)

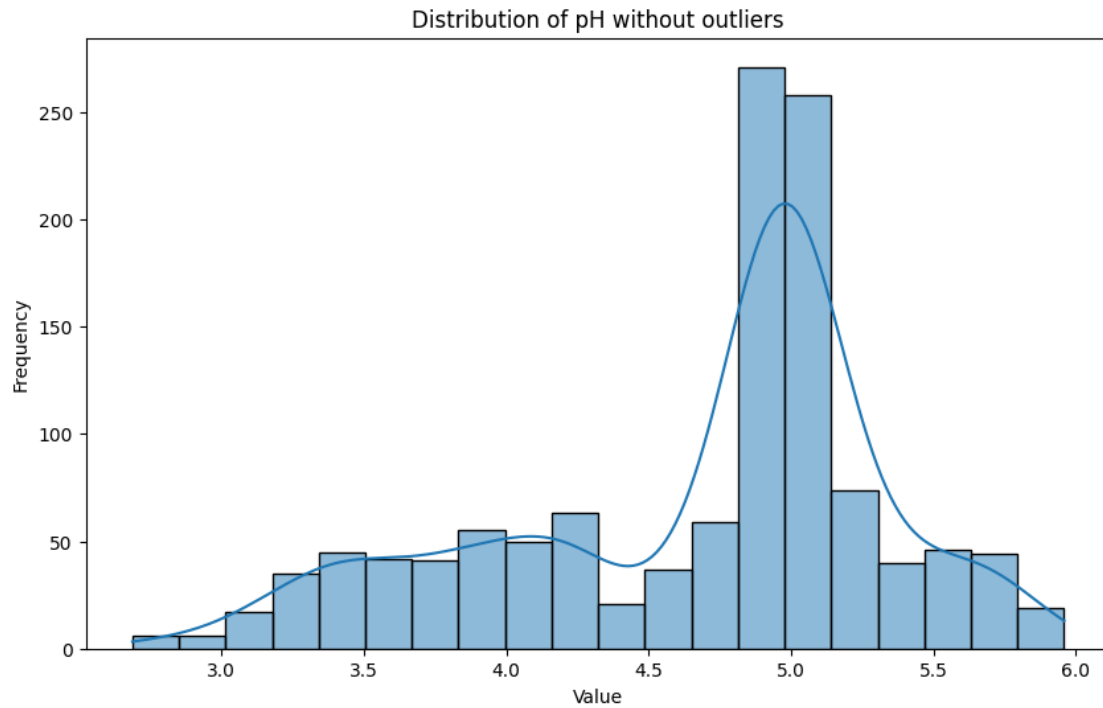
#Dropping rows with feature values that clearly looks like outliers.
# I did this manually by just looking at the plots, but i could've also removed
↳data with z-score
df_cleaned= df_cleaned.drop(df_cleaned[df_cleaned["Acoustic Firmness Index"] >
↳50].index)
df_cleaned= df_cleaned.drop(df_cleaned[df_cleaned["Odor index (a.u.)"] > 85].
↳index)

#I've could've also dropped these rows under, but the accuracy got worse when
↳doing this
#df_cleaned= df_cleaned.drop(df_cleaned[df_cleaned["Luminescence Intensity (a.u.
↳)"] > 0.017].index)

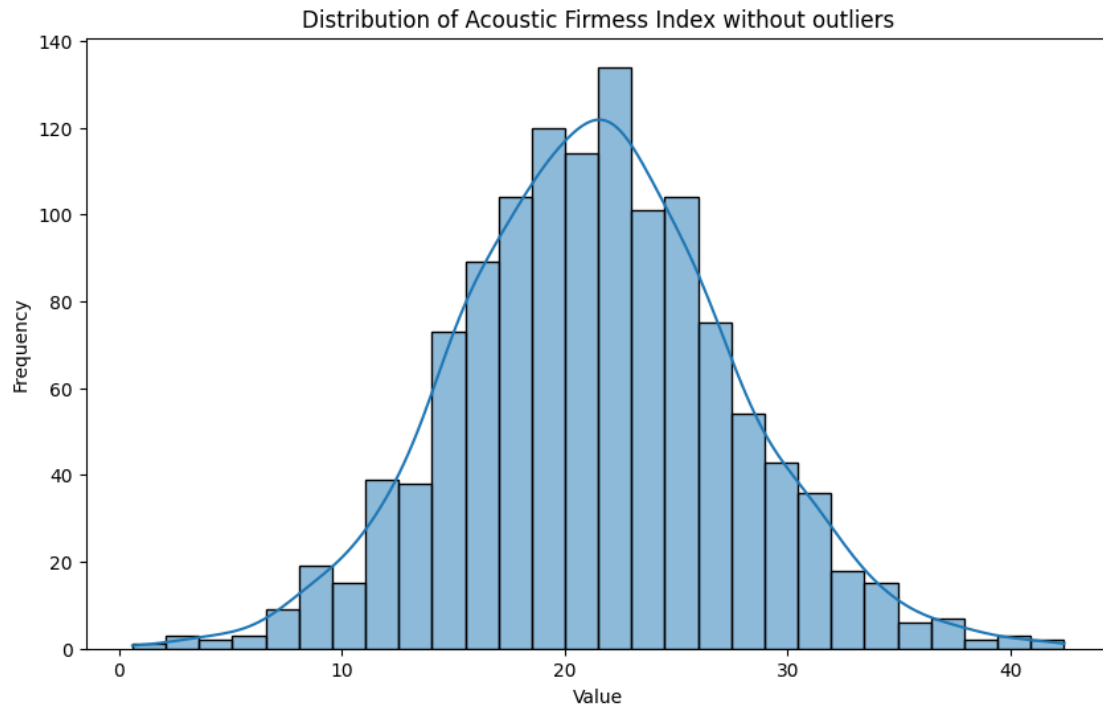
# For checking how many rows i've removed
#I just made comments out of this so you dont have to read all the output
#df_cleaned
```

0.0.5 Data preprocessing and visualisation

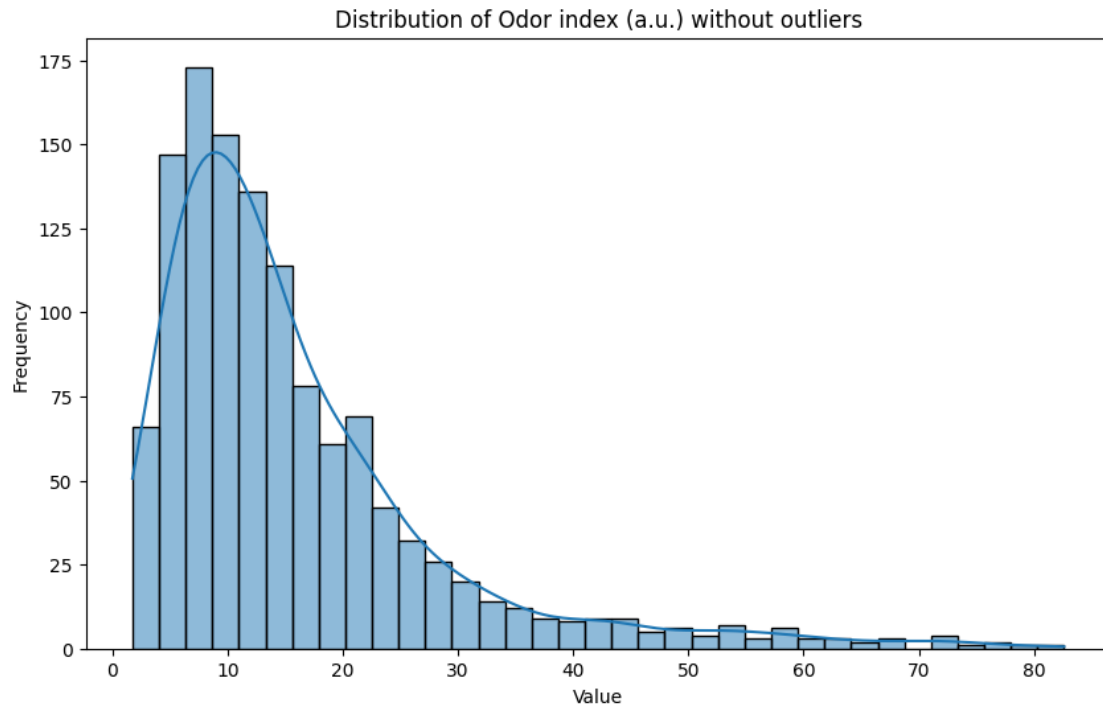
```
[1030]: # Plotting the distribution of pH values after data-cleaning
plt.figure(figsize=(10, 6))
sns.histplot(df_cleaned["pH"], kde=True)
plt.title('Distribution of pH without outliers')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```



```
[1031]: # Plotting the distribution of Acoustic Firmness Index after data-cleaning
plt.figure(figsize=(10, 6))
sns.histplot(df_cleaned["Acoustic Firmness Index"], kde=True)
plt.title('Distribution of Acoustic Firmness Index without outliers')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```



```
[1032]: # Plotting the distribution of Odor index (a.u.) after data-cleaning
plt.figure(figsize=(10, 6))
sns.histplot(df_cleaned["Odor index (a.u.)"], kde=True)
plt.title('Distribution of Odor index (a.u.) without outliers')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```



```
[1033]: # Removing the target column from the DataFrame and uses the rest as features
X = df_cleaned.drop('Edible', axis=1)
# Selcts Edible as target columns
y = df_cleaned['Edible']

#Splitting up into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42, stratify=y)

# Initialise standard scaler and compute mean and stddev from training data
sc = StandardScaler()
sc.fit(X_train)

# Transform (standardise) both X_train and X_test with mean and stddev from
# training data, to avoid leakage
X_train_sc = sc.transform(X_train)
X_test_sc = sc.transform(X_test)
```

0.0.6 Modelling

```
[1034]: # No need to standardize the data for Random Forest

# Trying ot RandomForestClassifier
rfc = RandomForestClassifier(random_state=42)
rfc.fit(X_train, y_train)

# Making predictions with training data
y_pred_train = rfc.predict(X_train)
#computing accuracy from training data
accuracy_train = accuracy_score(y_train, y_pred_train)
print(f'Accuracy on training data: {accuracy_train:.2f}')

# Making predictions with test data
y_pred = rfc.predict(X_test)
#computing accuracy from test data
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy on test data: {accuracy:.2f}')
```

Accuracy on training data: 1.00

Accuracy on test data: 0.92

Here i tried different classifiers, but i took it out so you have less code to look at!

0.0.7 Final evaluation

Got best accuracy from randomforest, so i use this classifier

```
[1041]: # Getting feature importance
feature_importances = rfc.feature_importances_

# Creating a DataFrame for feature importances
features_df = pd.DataFrame({'Feature': X.columns, 'Importance':_
    ↪feature_importances})

# Sorting the features by importance
features_df = features_df.sort_values(by='Importance', ascending=False)
print("Feature importances:\n", features_df)
```

Feature importances:

	Feature	Importance
15	pH	0.234173
13	Sugar Content (mg)	0.126118
2	Bitterness Scale	0.117119
14	Weight (mg)	0.080775
6	Length (mm)	0.069948
7	Luminescence Intensity (a.u.)	0.068749
9	Odor index (a.u.)	0.056343

3	Circumference (mm)	0.047983
0	Acoustic Firmness Index	0.038648
8	Magnetic orientation (degree)	0.026284
1	Atmospheric Pressure at Harvest (Pa)	0.025742
4	Color Intensity (a.u.)	0.025164
5	Find Distance from Main Vulcano (km)	0.024629
12	Soil pH where Grown	0.023813
10	Seed Count	0.018422
11	Skin Thickness (mm)	0.016090

In the next box i try all different combinations of features. I use the feautre combination that has the highest accuracy

```
[1042]: # Selecting subsets of features starting from the most important
best_average_accuracy = 0
best_features = []

for i in range(1, len(features_df.Feature) + 1):
    #Selecting the combination to use for the model
    selected_features = features_df.Feature.head(i)
    X_selected = X[selected_features]

    #storing accuracy
    accuracies = []

    # Iterate over a range of random states for train-test split
    for r in range(1,100, 10):
        # Splitting data with varying random states, to avoid overfitting
        #splitting with only the selected features
        X_train, X_test, y_train, y_test = train_test_split(X_selected, y,
↳test_size=0.2, random_state=r)

        # Initialize the model
        model = RandomForestClassifier(random_state=42)
        model.fit(X_train, y_train)

        # Evaluate the model on the test set
        predictions = model.predict(X_test)
        accuracy = accuracy_score(y_test, predictions)
        accuracies.append(accuracy)

    # Calculate the average accuracy across all splits for this set of features
    average_accuracy = np.mean(accuracies)

    # Update the best feature-lsit if the current average accuracy is better
    if average_accuracy > best_average_accuracy:
        best_average_accuracy = average_accuracy
```

```

best_features = selected_features.tolist()

print(f"Best Average Accuracy: {best_average_accuracy:.4f}")
print(f"Best Features: {best_features}")

```

Best Average Accuracy: 0.9106

Best Features: ['pH', 'Sugar Content (mg)', 'Bitterness Scale', 'Weight (mg)', 'Length (mm)', 'Luminescence Intensity (a.u.)', 'Odor index (a.u.)', 'Circumference (mm)', 'Acoustic Firmness Index']

In the next box, i try to find the best paramters for the model. I started with looping over `n_estimators`, and selected the value with the best accuracy. Then i adjusted the code with `best_n_estimators`, and looped over different values of `max_depth`. I select the `max_depth`-value that gave me the highest accuraccy with the `best_n_estimators`. This goes on in the same pattern for the next parameter.

I've tried using different criterions and `random_states` for the model, and earlier i've found that `criterion= "gini"` and `randomstate= 42` has worked well for me..

I could have tried to find more paramters to use in the model, but the ones i've used here has worked well enough.

```

[1051]: # Here i set the best values for parameters i've already found.
best_n_estimators = 181
best_max_depth= 11

# Stores the best configuration in the loop
best_accuracy = 0
best_config = {}

#looping over different values for the next parameter to find, which here is min_samples_split
parameters_to_try= [n for n in range(2,30,3)]

for parameter in parameters_to_try:
    #Storing accuracies for different random states
    accuracies = []

    # Iterate over a range of random states for train-test split
    for r in range(1, 100, 10):

        # Splitting data with varying random states, to avoid overfitting
        #also only using the selected features
        X_train, X_test, y_train, y_test = train_test_split(X[best_features], y,
        y, test_size=0.2, random_state=r)

        # Initialize and train the RandomForest model
        #Using the best paramters that i've found earlier with the same code

```

```

        #using different values for min_samples_split to find the new best one
        ↳ in this case
        model = RandomForestClassifier(random_state=42, criterion="gini",
        ↳ n_estimators=best_n_estimators, max_depth= best_max_depth,
        ↳ min_samples_split= parameter)
        model.fit(X_train, y_train)

        # Evaluate the model on the validation set
        predictions = model.predict(X_test)
        accuracy = accuracy_score(y_test, predictions)
        accuracies.append(accuracy)

        # Calculate the average accuracy across all splits
        average_accuracy = np.mean(accuracies)

        # Update the best configuration for new parameter if necessary
        if average_accuracy > best_accuracy:
            best_accuracy = average_accuracy
            best_parameter= parameter
            best_config = {
                'n_estimators': best_n_estimators,
                'max_depth': best_max_depth,
                'random_state': 42,
                'criterion': "gini",
                'min_samples_split': best_parameter #store best new parameter
                #if im trying to find more paramters:
                #'min_samples_split': best_mms
                #'new_parameter': best_parameter
            }

        # Print the best configuration and its accuracy
        print("Best Configuration:")
        print(best_config)
        print(f"Best Validation Accuracy: {best_accuracy:.4f}")

```

```

Best Configuration:
{'n_estimators': 181, 'max_depth': 11, 'random_state': 42, 'criterion': 'gini',
 'min_samples_split': 2}
Best Validation Accuracy: 0.9114

```

```

[1053]: # Just checking what parameters i ended up with
        best_config

```

```

[1053]: {'n_estimators': 181,
        'max_depth': 11,
        'random_state': 42,
        'criterion': 'gini',

```

```
'min_samples_split': 2}
```

```
[1054]: # Create a new RandomForestClassifier with the best configuration
final_model= RandomForestClassifier(**best_config)

#Training model with the whole dataset, but only with the selected features
final_model.fit(X[best_features], y)

#loading up the test-set
x_val= pd.read_csv("test.csv",index_col=0 )
#Selecting only the best feautres
x_val_selected = x_val[best_features]

#predicting
predictions= final_model.predict(x_val_selected).astype(int)
```

```
[1055]: # Create a DataFrame with predictions
predictions_df = pd.DataFrame(predictions, columns=['Edible'])
predictions_df.index.name = 'index'

# Save to a CSV file
predictions_df.to_csv('predictions.csv')
```